Cue — AI Pipeline (proposal view, not an implemented system)
Inference path (product) and training/evaluation path (research, independent of the product database)

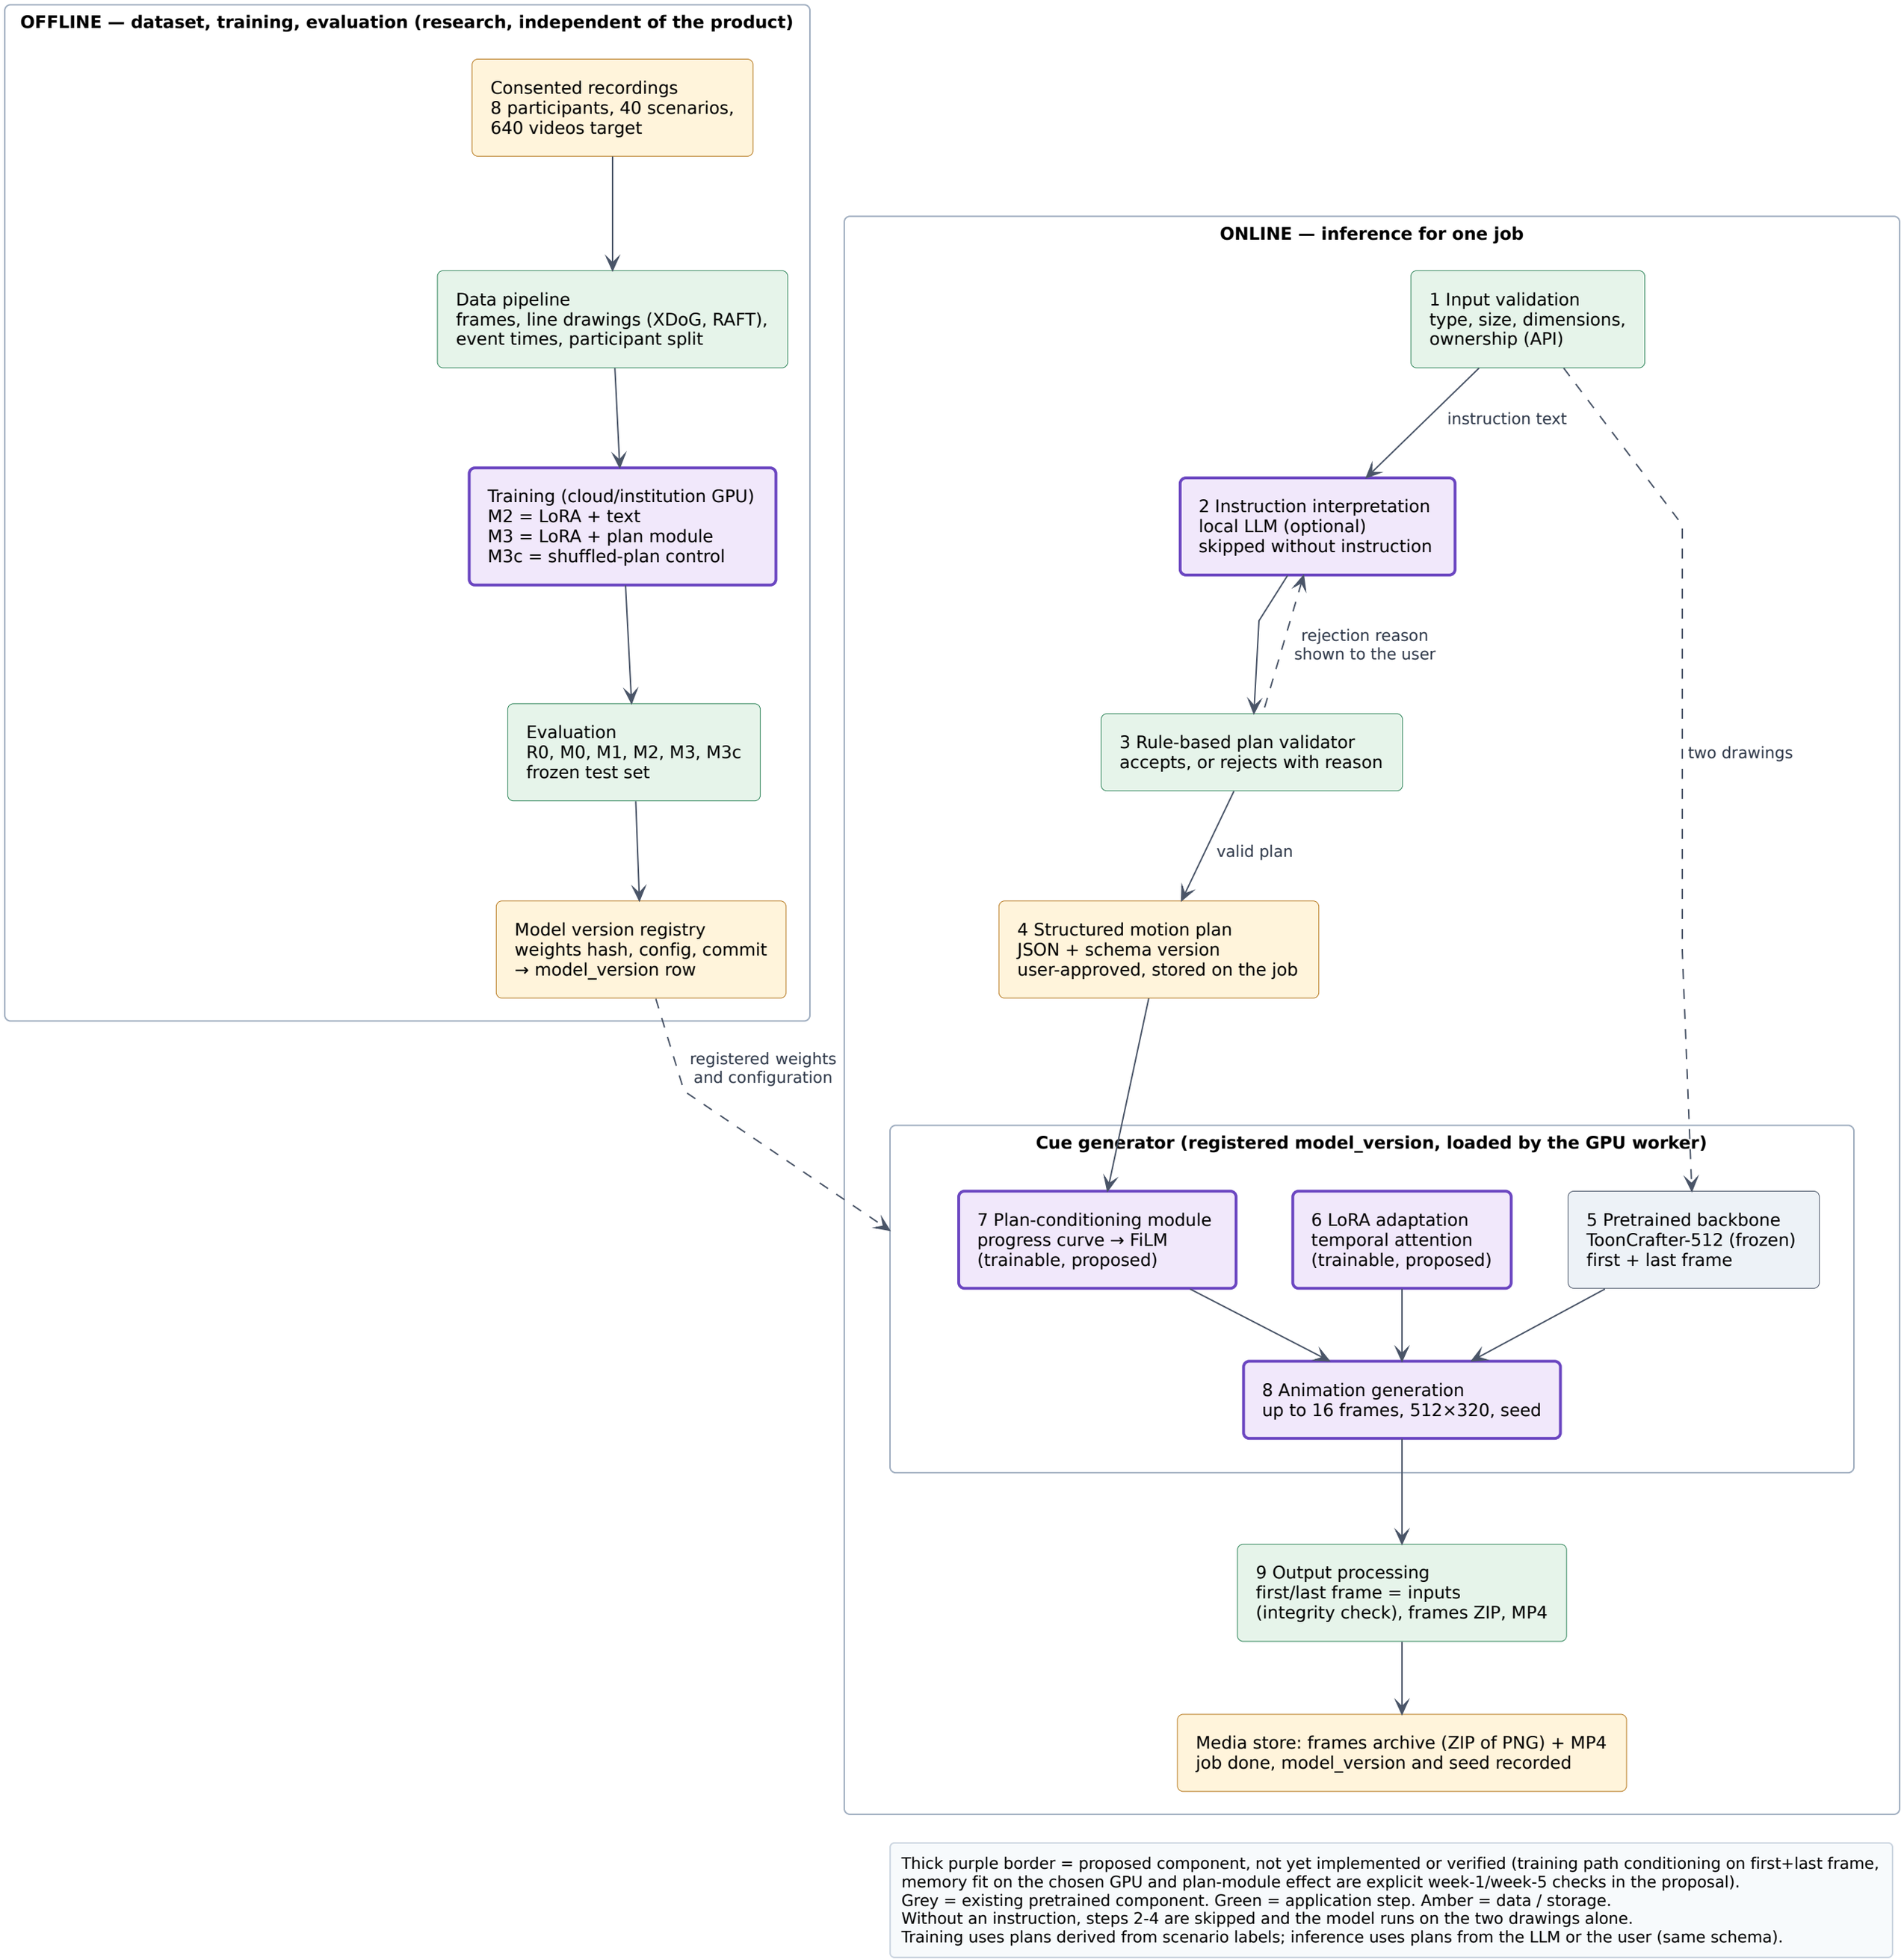

Thick purple border = proposed component, not yet implemented or verified (training path conditioning on first+last frame, memory fit on the chosen GPU and plan-module effect are explicit week-1/week-5 checks in the proposal).
Grey = existing pretrained component. Green = application step. Amber = data / storage.
Without an instruction, steps 2-4 are skipped and the model runs on the two drawings alone.
Training uses plans derived from scenario labels; inference uses plans from the LLM or the user (same schema).